

```

#!/usr/bin/env python3
# -*- coding: utf-8 -*-
"""
Created on Mon May  9 11:59:16 2022

@author: hillarywolff with Jason Winik
"""
import pandas as pd
import numpy as np
from numpy import mean
from numpy import absolute
from numpy import arange
import random

from sklearn.model_selection import train_test_split, cross_val_score,
LeaveOneOut, RepeatedKFold
from sklearn.linear_model import LogisticRegression, LinearRegression,
Lasso, LassoCV, Ridge, RidgeCV
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import PolynomialFeatures

import statsmodels.api as sm
import statsmodels.formula.api as smf

import matplotlib.pyplot as plt
import seaborn as sns

from itertools import combinations

# to ignore warnings
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)

# 1. Chapter 5
# a. *question 6*
random.seed(5)
# 6. We continue to consider the use of a logistic regression model to
# predict the probability of default using income and balance on the
# Default data set. In particular, we will now compute estimates for
# the standard errors of the income and balance logistic regression
# coefficients in two different ways: (1) using the bootstrap, and (2)
# using
# the standard formula for computing the standard errors in the glm()
# function. Do not forget to set a random seed before beginning your
# analysis.
PATH = r"/Users/hillarywolff/Documents/GitHub/machine_learning/PS3/"
df = pd.read_csv(PATH + 'Default.csv')
# (a) Using the summary() and glm() functions, determine the estimated
# standard

```

```
# errors for the coefficients associated with income and balance in a
multiple
# logistic regression model that uses both predictors.
df = df.drop('Unnamed: 0', axis=1)

df['default'] = np.where((df['default'].str.contains('Yes')), 1, 0)
df['student'] = np.where((df['student'].str.contains('Yes')), 1, 0)

X = df[['balance', 'income']]
X = sm.add_constant(X)
y = df['default']

results = sm.Logit(y, X).fit().summary()
print(results)
```

```
#                               Logit Regression Results
#
=====
# Dep. Variable:                default    No. Observations:
10000
# Model:                        Logit      Df Residuals:
9997
# Method:                       MLE       Df Model:
2
# Date:                         Mon, 09 May 2022    Pseudo R-squ.:
0.4594
# Time:                         16:05:26    Log-Likelihood:
-789.48
# converged:                    True       LL-Null:
-1460.3
# Covariance Type:             nonrobust    LLR p-value:
4.541e-292
#
=====
#                               coef    std err          z      P>|z|      [0.025
0.975]
#
-----
# const          -11.5405         0.435    -26.544     0.000    -12.393
-10.688
# balance         0.0056         0.000     24.835     0.000     0.005
0.006
# income         2.081e-05    4.99e-06     4.174     0.000    1.1e-05
3.06e-05
#
=====
=====
```

```

# std err for balance = 0.00, std err for income = 4.99e-6

# (b) Write a function, boot.fn(), that takes as input the Default
data
# set as well as an index of the observations, and that outputs
# the coefficient estimates for income and balance in the multiple
# logistic regression model.

def get_indices(data, num_samples):
    """
    Gets a random subsample (based on num_samples) of the
    indexes of the dataset (data)
    """

    return np.random.choice(data.index, int(num_samples),
                             replace=True)

def boot_fn(data, index):
    """
    Runs one logistic regression on only the indices specified
    on index that are found in data. It then returns the three
    coefficients associated with the regression.
    """
    X = data[['balance', 'income']].loc[index]
    X = sm.add_constant(X)
    y = data['default'].loc[index]

    lr = sm.Logit(y, X).fit(dis=0)
    intercept = lr.params[0]
    coef_balance = lr.params[1]
    coef_income = lr.params[2]
    return [intercept, coef_balance, coef_income]

# (c) Use the boot() function together with your boot.fn() function to
# estimate the standard errors of the logistic regression coefficients
# for income and balance.
def boot(data, func, R):
    intercept = []
    coeff_balance = []
    coeff_income = []
    for i in range(R):

        [inter, balance, income] =
func(data, get_indices(data, len(data)))
        intercept.append(float(inter))
        coeff_balance.append(balance)

```

```

    coeff_income.append(income)

    intercept_statistics =
{'estimated_value':np.mean(intercept),'std_error':np.std(intercept)}
    balance_statistics =
{'estimated_value':np.mean(coeff_balance),'std_error':np.std(coeff_bal
ance)}
    income_statistics =
{'estimated_value':np.mean(coeff_income),'std_error':np.std(coeff_inco
me)}
    return
{'intercept':intercept_statistics,'balance_statistics':balance_statist
ics,'income_statistics':income_statistics}

results = boot(df, boot_fn, 1000)
print('Intercept - ', results['intercept'])
print('Balance - ', results['balance_statistics'])
print('Income - ', results['income_statistics'])

# (d) Comment on the estimated standard errors obtained using the
# glm() function and using your bootstrap function.

# the standard errors from the Logit function were incredibly similar
to the
# standard errors in the bootstrapping function.

# std err balance: 0.0002, std err income: 4.87e-6

#####

# b. *question 8*

# Generate a simulated data set as follows:
# > set.seed (1)
# > x <- rnorm (100)
# > y <- x - 2 * x^2 + rnorm (100)
# In this data set, what is n and what is p? Write out the model
# used to generate the data in equation form.

sim_df = pd.DataFrame()
N = 100
sim_df['x'] = np.random.normal(0, 1, N)
sim_df['y'] = sim_df['x']-2 * sim_df['x'].pow(2) +
np.random.normal(0,1,100)

# N = 100 and p = 2 which is found by looking at  $Y = X - 2X^2 + e$ 

```

```

# (b) Create a scatterplot of X against Y . Comment on what you find.

sns.scatterplot(sim_df['x'], sim_df['y'])

# The data is quadratic which we know from the exponent, but a
# majority of
# points are located between -1 and 1 and minics a normal distribution
# which is
# expected since the simulated dataframe used a mean of 0 and stdev of
# 1 to
# generate the points.

# (c) Set a random seed, and then compute the L00CV errors that
# result from fitting the following four models using least squares:
# i.  $Y = \beta_0 + \beta_1 X + \epsilon$ 
# ii.  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \epsilon$ 
# iii.  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \epsilon$ 
# iv.  $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \beta_4 X^4 + \epsilon$ 

# Note you may find it helpful to use the data.frame() function
# to create a single data set containing both X and Y .
random.seed(20)
sim_df['x2'] = np.power(sim_df['x'], 2)
sim_df['x3'] = np.power(sim_df['x'], 3)
sim_df['x4'] = np.power(sim_df['x'], 4)

X1 = sim_df[['x']]
X2 = sim_df[['x', 'x2']]
X3 = sim_df[['x', 'x2', 'x3']]
X4 = sim_df[['x', 'x2', 'x3', 'x4']]
y = sim_df['y']

cv = LeaveOneOut()
model = LinearRegression()

cols = [X1, X2, X3, X4]

def MSE_L00CV(cols):
    mse_list = []
    for col in cols:
        scores = cross_val_score(model, col, y,
                                scoring='neg_mean_squared_error',
                                cv=cv)
        mse_list.append(mean((absolute(scores))))

    return mse_list

sim_mse = MSE_L00CV(cols)
print(sim_mse)

```

```

# [i: 4.757748722926266, ii: 1.6391419532039742, iii:
1.6672891086484705,
# iv: 1.6796562364607772]

# (d) Repeat (c) using another random seed, and report your results.
# Are your results the same as what you got in (c)? Why?

random.seed(54)
cols = [X1, X2, X3, X4]
new_mse = MSE_LOOCV(cols)
print(new_mse)

# [i: 4.757748722926266, ii: 1.6391419532039742, iii:
1.6672891086484705,
# iv: 1.6796562364607772]

# the results are the same because LOOCV uses N folds from the same
dataset, so
# any iteration of it will be the same.

# (e) Which of the models in (c) had the smallest LOOCV error? Is
# this what you expected? Explain your answer.

# the quadratic model had the smallest LOOCV error. this is expected
since we
# saw in our scatterplot that there was a quadratic relationship.

# (f) Comment on the statistical significance of the coefficient
estimates that
# results from fitting each of the models in (c) using least squares.
Do these
# results agree with the conclusions drawn based on the cross-
validation results?
result = smf.ols(formula="y ~ x+x2", data=sim_df).fit().summary()
print(result)

result = smf.ols(formula="y ~ x+x2+x3", data=sim_df).fit().summary()
print(result)

result = smf.ols(formula="y ~ x+x2+x3+x4",
data=sim_df).fit().summary()
print(result)

#                                     OLS Regression Results
#
=====
=====
# Dep. Variable:                    y    R-squared:

```

```

0.687
# Model:                                OLS    Adj. R-squared:
0.674
# Method:                               Least Squares    F-statistic:
52.13
# Date:                                Mon, 09 May 2022    Prob (F-statistic):
3.67e-23
# Time:                                17:38:51    Log-Likelihood:
-161.66
# No. Observations:                    100    AIC:
333.3
# Df Residuals:                        95    BIC:
346.3
# Df Model:                            4
# Covariance Type:                    nonrobust
#
=====
=====
#           coef      std err          t      P>|t|      [0.025
0.975]
#
-----
# Intercept      0.0714      0.200      0.358      0.721     -0.325
0.468
# x              1.1071      0.279      3.966      0.000      0.553
1.661
# x2             -2.5017      0.372     -6.732      0.000     -3.239
-1.764
# x3             -0.0546      0.152     -0.359      0.721     -0.357
0.248
# x4              0.1792      0.116      1.548      0.125     -0.051
0.409
#
=====
=====
# Omnibus:                3.559    Durbin-Watson:
1.999
# Prob(Omnibus):          0.169    Jarque-Bera (JB):
2.687
# Skew:                   0.259    Prob(JB):
0.261
# Kurtosis:               2.386    Cond. No.
14.0
#
=====
=====

# Notes:
# [1] Standard Errors assume that the covariance matrix of the errors

```

is correctly specified.

```
# the results from our OLS regression are in line with our L00CV
results where
# our x and x2 models are significant while x3 and x4 are not.
```

```
#####
#####
```

```
# 2.
# a. *question 11*
# We will now try to predict per capita crime rate in the Boston data
# set.
df = pd.read_csv(PATH+'Boston.csv')
```

```
# (a) Try out some of the regression methods explored in this chapter,
# such as best subset selection, the lasso, ridge regression, and
# PCR. Present and discuss results for the approaches that you
# consider.
```

```
boston = df
# forward stepwise:
```

```
#Add constant to dataframe
boston['constant'] = 1
#specify target
Y = boston['CRIM']
#Variables to use in forward propagation
vars_left_add = boston.columns.tolist()
vars_left_add = [e for e in vars_left_add if e not in ('CRIM',
'constant')]
#Regression type
ols = LinearRegression()
#Starting variables (only constant)
current_vars = ['constant']

X = boston[current_vars]
benchmark_error = np.mean(-1*cross_val_score(ols, X, Y, cv = 5,
scoring = 'neg_mean_squared_error'))
print(' Initial run with only one var (constant term/only bias
weight):', current_vars)
print('      Benchmark error:', benchmark_error)
print('')
```

```
for iter in range(len(vars_left_add)):
    print('\033[1m'+ 'Iteration:', iter, '\033[0m')
    error_list = []
    for var in vars_left_add: #For each variable that we can add
        #Modify X according to current iteration
```



```

# Dep. Variable:          Y    R-squared:
0.418
# Model:                  OLS   Adj. R-squared:
0.415
# Method:                 Least Squares   F-statistic:
120.3
# Date:                   Mon, 09 May 2022   Prob (F-statistic):
1.00e-58
# Time:                   19:57:43   Log-Likelihood:
-1669.0
# No. Observations:      506   AIC:
3346.
# Df Residuals:          502   BIC:
3363.
# Df Model:              3
# Covariance Type:       nonrobust
#

```

```

=====
=====
#               coef      std err          t      P>|t|      [0.025
0.975]
#
-----
# Intercept      -2.4701      0.362      -6.815      0.000      -3.182
-1.758
# constant      -2.4701      0.362      -6.815      0.000      -3.182
-1.758
# RAD           0.5281      0.039      13.578      0.000      0.452
0.605
# LSTAT         0.2574      0.049       5.203      0.000      0.160
0.355
# ZN            0.0205      0.014       1.476      0.140      -0.007
0.048
#

```

```

=====
=====
# Omnibus:          676.740   Durbin-Watson:
1.459
# Prob(Omnibus):    0.000   Jarque-Bera (JB):
88649.095
# Skew:            6.798   Prob(JB):
0.00
# Kurtosis:        66.402   Cond. No.
1.85e+16
#
=====
=====

```

```

np.random.seed(5)

```

```

X = boston[current_vars]
Y = boston['CRIM']
ols = LinearRegression()
np.mean(-1*cross_val_score(ols, X, Y, cv = 5, scoring =
'neg_mean_squared_error'))
# MSE = 44.46

# backward stepwise
vars_left_to_drop = boston.columns.tolist()
vars_left_to_drop = [e for e in vars_left_to_drop if e not in ('CRIM',
'constant')]
#Regression type
ols = LinearRegression()
#Starting variables (only constant)
current_vars = ['constant'] + vars_left_to_drop

X = boston[current_vars]
benchmark_error = np.mean(-1*cross_val_score(ols, X, Y, cv = 5,
scoring = 'neg_mean_squared_error'))
print(' Initial run with all vars:', current_vars)
print('      Benchmark error:', benchmark_error)
print('')

for iter in range(len(vars_left_to_drop)):
    print('\033[1m'+ 'Iteration:', iter, '\033[0m')
    error_list = []
    for var in vars_left_to_drop: #For each variable that we can add
        #Modify X according to current iteration
        vars_to_be_used = ['constant'] + [i for i in vars_left_to_drop
if i != var]
        X = boston[['constant'] + [i for i in vars_left_to_drop if i !
= var]]
        #Perform 5-fold CV to get errors
        error = np.mean(-1*cross_val_score(ols, X, Y, cv = 5, scoring
= 'neg_mean_squared_error'))
        error_list.append(error)
        print(' Running model with:', vars_to_be_used)
        print('      Error:', error)

    # Chose the smallest error
    min_error = min(error_list)
    chosen_col_index = error_list.index(min_error)

    # If our current smallest error is smaller than our previous
error, than we drop the variable associated with it
    # if not, we keep our model

```

```

        if min_error<benchmark_error:
            print('          *** Will drop:',
vars_left_to_drop[chosen_col_index])
            print('          *** Min error selected:', min_error)
            print('          *** Chose the variable that generated the min
error + was lower than previous error')
            print('')
            # Add the variable that produced the smallest error to
current_vars
            current_vars = vars_to_be_used
            del vars_left_to_drop[chosen_col_index] #delete chosen
variable from vars_left_to_drop
            benchmark_error = min_error # Update benchmark_error
        else:
            print('          \033[4m*** No variable was selected',
'\033[0m')
            print('          *** Previous error rate (',
benchmark_error,') is lower than smallest error rate of this iteration
(', min_error ,')')
            print('          *** Break')
            break

```

```

print('')
print('Variables chosen for our model', current_vars)

```

```

# Variables chosen for our model ['constant', 'ZN', 'INDUS', 'DIS',
'RAD', 'TAX', 'PTRATIO', 'LSTAT'] with error 44.57
result = smf.ols(formula="Y ~
constant+RAD+LSTAT+ZN+INDUS+DIS+PTRATIO+TAX",
data=boston).fit().summary()

```

```

#                               OLS Regression Results
#

```

```

=====
=====
# Dep. Variable:                Y    R-squared:
0.425
# Model:                        OLS    Adj. R-squared:
0.417
# Method:                       Least Squares    F-statistic:
52.64
# Date:                         Mon, 09 May 2022    Prob (F-statistic):
4.56e-56
# Time:                         19:59:32    Log-Likelihood:
-1666.0
# No. Observations:             506    AIC:
3348.
# Df Residuals:                 498    BIC:
3382.
# Df Model:                     7

```

```

# Covariance Type:          nonrobust
#
=====
=====
#               coef      std err          t      P>|t|      [0.025
0.975]
#
-----
# Intercept      -0.9394      1.510      -0.622      0.534      -3.907
2.028
# constant      -0.9394      1.510      -0.622      0.534      -3.907
2.028
# RAD           0.5326      0.085       6.232      0.000       0.365
0.700
# LSTAT         0.2641      0.053       4.939      0.000       0.159
0.369
# ZN            0.0372      0.019       2.005      0.045       0.001
0.074
# INDUS        -0.1007      0.081      -1.250      0.212      -0.259
0.058
# DIS          -0.5403      0.236      -2.289      0.022      -1.004
-0.077
# PTRATIO       0.0025      0.167       0.015      0.988      -0.327
0.332
# TAX          -0.0006      0.005      -0.121      0.904      -0.011
0.009
#
=====
=====
# Omnibus:                678.269   Durbin-Watson:
1.484
# Prob(Omnibus):          0.000   Jarque-Bera (JB):
89641.511
# Skew:                   6.823   Prob(JB):
0.00
# Kurtosis:               66.762   Cond. No.
4.57e+20
#
=====
=====

np.random.seed(5)

X = boston[current_vars]
Y = boston['CRIM']
ols = LinearRegression()
np.mean(-1*cross_val_score(ols, X, Y, cv = 5, scoring =
'neg_mean_squared_error'))
# MSE= 45

```

(b) Propose a model (or set of models) that seem to perform well on
this data set, and justify your answer. Make sure that you are
evaluating
model performance using validation set error, crossvalidation, or
some other
reasonable alternative, as opposed to using training error.

our forward stepwise method produced the lowest error which was
44.46 with
variables ['constant', 'RAD', 'LSTAT', 'ZN']

(c) Does your chosen model involve all of the features in the data
set? Why or why not?

no, it only involves those most pertinent variables that are
relevant for the
analysis. The process of stepwise goes through all possible
combinations
of variables to find the lowest error rate and therefore the best
predictors

